

Combinational Building Blocks

Index


- Building Block definition
- Multiplexer
- Decoders
- Adders
- Subtractor
- Comparators
- Arithmetic/Logical Unit (ALU)
- Shifters and Rotators
- Multiplier and divider
- Floating-Point Addition

Building Block Definition



- Combinational logic can be grouped in **larger building blocks** to build more complex systems
- **Abstraction, hierarchy, modularity, and regularity**
 - **abstract** (hide) gate-level details to **emphasize the function** of the **module** (block)
 - **hierarchically** assembled from simpler components
 - **well-defined interfaces** in order to use them as a black boxes
- Examples:
 - seven-segment display decoders (already considered)
 - multiplexers
 - decoders
 - arithmetic circuits
- We will use many of these building blocks to build a **microprocessor**

Multiplexer (mux)

- A device that **choose an output among several possible inputs**, based on the value of a **select signal**
- A 2:1 multiplexer has two data inputs D and one output Y and a select signals S:
 - if $S=0$, $Y=D_0$
 - if $S=1$, $Y=D_1$






S	D ₀	D ₁	Y
0	0	0	0
0	0	1	0
0	1	0	1
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	0
1	1	1	1



Wider multiplexer


- A 4:1 multiplexer has four data inputs and one output and **two select signals** are needed to choose among the four data inputs
 - it can be built using sum-of-product logic, or multiple 2:1 multiplexers



- Wider multiplexers (8:1 and 16:1) can be built by expanding the method
 - in general, a **N:1 multiplexer needs $\log_2 N$ select lines**


Multiplexer example

- A multiplexer is like a **data router**: connects only one input at a time, **allowing different data sources to share the same output line efficiently**
- This makes it possible to **handle many inputs in a controlled, sequential way** without needing separate circuits for each one



Multiplexer Logic (1)


- Multiplexers can be used as **lookup tables** to **perform logic functions**
- For example, we can implement a two-input AND gate:



- A and B serve as select lines
- inputs are connected to 0 or 1, **according to the truth table**
- changing the data inputs, the multiplexer **can be reprogrammed** to perform a different function
- In general, a **2^N-input multiplexer can be programmed to perform any N-input logic function** by applying 0 and 1 to the appropriate data inputs


Multiplexer Logic (2)

- We can **cut the multiplexer size in half**, using only a 2^{N-1} input multiplexer to perform any N-input logic function
- Provide **one of the literals to the multiplexer inputs**:




Multiplexer Logic (3)

- As an example, implement the following function $Y = A \bar{B} + \bar{B} \bar{C} + \bar{A} B C$
 - using a 8:1 multiplexer:



- and a 4:1 multiplexer:



Multiplexer VHDL (1)

```
entity mux_4_1 is
    port( s0, s1: in std_logic;
          a, b, c, d: in std_logic;
          y: out std_logic
    );
end mux_4_1;

architecture rtl of mux_4_1 is
begin
    y <= a when s0='0' and s1='0' else
        b when s0='0' and s1='1' else
        c when s0='1' and s1='0' else
        d when s0='1' and s1='1' else
        '-';
end rtl;
```

Multiplexer VHDL (2)

```
entity mux_4_1_8bit is
    port( s0, s1: in std_logic;
          a, b, c, d: in std_logic_vector(0 to 7);
          y: out std_logic_vector(0 to 7)
    );
end mux_4_1_8bit;

architecture rtl of mux_4_1_8bit is
    signal sel: std_logic_vector(0 to 1);
begin
    with sel select
        y <= a when "00",
                  b when "01",
                  c when "10",
                  d when "11",
                  "-----" when others;
end rtl;
```


Multiplexer VHDL (3)

```
entity mux_4_1_Nbit is
    generic( N: integer );
    port( sel: in std_logic_vector(0 to 1);
          a, b, c, d: in std_logic_vector(0 to N-1);
          y: out std_logic_vector(0 to N-1)
        );
end mux_4_1_Nbit;
```


```
architecture rtl of mux_4_1_Nbit is
begin
    with sel select
        y <= a when "00",
                  b when "01",
                  c when "10",
                  d when "11",
                  (others => '-') when others;
end rtl;
```

Decoder

- A device with **N inputs** and **2^N outputs** that asserts **exactly one** of its outputs depending on the input combination
- For example, consider a 2:4 decoder
 - when $A_{1:0} = 00$, Y_0 is 1
 - when $A_{1:0} = 01$, Y_1 is 1
 - and so forth
- The outputs are called **one-hot**, because just one is “hot” (HIGH) at a given time
- In the implementation, each gate depends on either the true or the complementary form of each input
 - an **$N:2^N$ decoder** can be constructed from **2^N N-input AND gates** that accept the various combinations of true or complementary inputs
 - each output in a decoder represents a single minterm
 - for example, Y_3 represents the minterm $A_1 A_0$.



A_1	A_0	Y_3	Y_2	Y_1	Y_0
0	0	0	0	0	1
0	1	0	0	1	0
1	0	0	1	0	0
1	1	1	0	0	0




Decoder VHDL

```
entity decoder_2_4 is
    port( dec_in: in std_logic_vector(0 to 1);
          |   dec_out: out std_logic_vector(0 to 3)
        );
end decoder_2_4;

architecture rtl of decoder_2_4 is
begin
    with dec_in select
        dec_out <= "0001" when "00",
                    "0010" when "01",
                    "0100" when "10",
                    "1000" when "11",
                    "----" when others;
end rtl;
```


Decoder example

- A decoder can be used to control which of several lines is “on” based on the binary value of the input. It can be used in **memory address selection**:



Decoder Logic

- Decoders can be combined with OR gates **to build logic functions**
 - because each output of a decoder represents a single minterm, the function is built as the OR of all of the minterms in the function
 - example: a two-input XNOR function using a 2:4 decoder and a single OR gate




- An N -input function with M 1s in the truth table can be built with an $N:2N$ decoder and an M -input OR.

Half and Full Adders

- Addition is one of the most common operations in digital systems

- **Half adder**


- two inputs, A and B, and two outputs, S and C_{out}
- S is the sum of A and B
- C_{out} is the eventual carry out




- In a multi-bit adder, C_{out} is added (carried) in to the next most significant bit
 - half adder lacks a C_{in} input to accept the C_{out} of the previous column

- **Full Adder**


- accepts also the carry in C_{in}



Carry Propagate Adder



- A **N-bit adder** that sums two N-bit inputs and a carry-in to produce an N-bit result and a carry-out
 - the carry bit propagates into the adder
 - drawn like a full adder except that input/output are **busses** rather than single bits
- Simplest way to build: chain together N full adders (**Ripple-Carry Adder RCA**)
 - C_{ouy} of one stage acts as the C_{in} of the next stage



- A good application of modularity and regularity
 - full adder module is reused many times to form a larger system
- It is **slow when N is large**.
 - S₃ depends on C₂, which depends on C₁, and so forth
 - the delay of the adder t_{ripple} = N * t_{FA} grows directly with the number of bits


CPA VHDL

```
entity adder is
    generic(N: integer := 8);
    port(A, B: in std_logic_vector(N-1 downto 0);
         C_in: in std_logic;
         S: out std_logic_vector(N-1 downto 0);
         C_out: out std_logic);
end;


architecture synth of adder is
    signal result: std_logic_vector(N downto 0);
begin
    result <= ("0" & A) + ("0" & B) + C_in;
    S <= result(N-1 downto 0);
    Cout <= result(N);
end;
```

Carry-Lookahead Adder (1)

- RCA are slow because the carry must propagate:




- The idea is to **divide the adder into blocks** and determine the carry out of a block **as soon as** the carry in is known
 - it looks ahead across the blocks rather than wait to ripple through all the full adders



Carry-Lookahead Adder (2)


- **Carry-Lookahead Adder CLA** use **generate (G)** and **propagate (P)** signals
- A column “**generates**” a carry if it produces a carry independent of its carry-in
 - both input are 1 -> $G_i = A_i \text{ AND } B_i$
- A column “**propagates**” a carry if it produces a carry whenever there is a carry-in
 - either one of the input is 1 -> $P_i = A_i \text{ OR } B_i$



Carry-Lookahead Adder (2)


- We can extend to multiple-bit blocks
 - a block “generates” a carry if it produces a carry out independent of the carry-in
 - a block “propagate” a carry if it produces a carry out whenever there is a carry-in
- A block $j:i$ (spanning columns i through j) **generates** a carry
 - if the most significant column generates a carry
 - if the previous column generated a carry and the most significant column propagates it
 - and so forth
 - $G_{j:i} = G_j + P_j(G_{j-1} + P_{j-1}(G_{j-2} + P_{j-2}(\dots(G_{i+2} + P_{i+2}(G_{i+1} + P_{i+1}G_i)))$
 - e.g. $G_{7:4} = G_7 + P_7(G_6 + P_6(G_5 + P_5G_4))$
- A block $i:j$ **propagates** a carry
 - if all the columns in the block propagate the carry
 - $P_{j:i} = P_jP_{j-1}P_{j-2}\dots P_{i+2}P_{i+1}P_i$
 - e.g. $P_{3:0} = P_3P_2P_1P_0$
- So, we can quickly compute the carry out of the block
 - $C_j = G_{j:i} + P_{j:i}C_{i-1}$

Carry-Lookahead Adder (3)



Carry-Lookahead Adder (4)


- A 32-bit carry-lookahead adder composed of eight 4-bit blocks:



- Each block contains a 4-bit ripple-carry adder and the **lookahead** circuit to compute the carry out of the block given the carry-in
- All blocks compute column generate and propagate signals **simultaneously**
- Then C_{in} advances to C_{out} through each block until the last
 - C_{in} proceeds through the lookahead gates to produce C_3
 - C_3 proceeds through its lookahead block to produce C_7
 - C_7 proceeds through its lookahead block to produce C_{11}
 - and so on until C_{27} , the carry-in to the last block
- The last block contains a short ripple-carry adder (no more the lookahead logic)

Carry-Lookahead Adder delay

- An N-bit adder divided into k-bit blocks has a delay
 - t_{pg} is the delay to generate signals P_i and G_i
 - t_{pg_block} is the delay to generate signals $P_{j:i}$ and $G_{j:i}$ for a k-bit block
 - $t_{lookahead}$ is the delay from C_{in} to C_{out} through the logic of the k-bit CLA block



- $t_{CLA} = t_{pg} + t_{pg_block} + (N/k - 1) t_{lookahead} + k t_{FA}$
 - For $N > 16$, the carry-lookahead adder is much faster than the ripple-carry adder
 - however, the adder delay still increases linearly with N
- **Faster adders: more hardware, more expensive and power-hungry**
 - **trade-offs** must be considered when choosing an appropriate adder for a design

Ripple-Carry Adder vs Carry-Lookahead Adder

- Compare the delays of a 32-bit ripple-carry adder and a 32-bit carry-lookahead adder with 4-bit blocks
- Assume that each gate delay is 100ps and that a full adder delay is 300ps
- The propagation delay of the 32-bit ripple-carry adder is
 - $t_{\text{ripple}} = N * T_{\text{FA}} = 32 * 300\text{ps} = 9.6 \text{ ns}$
- The carry-lookahead adder has
 - $t_{\text{pg}} = 100\text{ps}$ (one AND or one OR)
 - $t_{\text{pg_block}} = 6 * 100 \text{ ps} = 600\text{ps}$ (six AND/OR gates)
 - $t_{\text{lookahead}} = 2 * 100 \text{ ps} = 200\text{ps}$ (one OR and one AND)
 - $t_{\text{CLA}} = t_{\text{pg}} + t_{\text{pg_block}} + (N/k-1)t_{\text{lookahead}} + kt_{\text{FA}} = 100 + 600 + (32/4 - 1)*200 + 4*300 = 3.3\text{ns}$
- The carry-lookahead adder is almost 3 times faster than the ripple-carry adder

CLA VHDL

```
entity CLA4 is
    port (A, B      : in  std_logic_vector(3 downto 0);
          C_in     : in  std_logic;
          S        : out std_logic_vector(3 downto 0);
          C_out    : out std_logic);
end CLA4;

architecture behavioral of CLA4 is
    signal G, P : std_logic_vector(3 downto 0);
    signal C     : std_logic_vector(4 downto 0);

begin
    -- Initial carry
    C(0) <= C_in;

    -- Generate and Propagate
    G <= A and B;
    P <= A xor B;


    -- Carry Lookahead Logic
    C(1) <= G(0) or (P(0) and C(0));
    C(2) <= G(1) or (P(1) and G(0)) or (P(1) and P(0) and C(0));
    C(3) <= G(2) or (P(2) and G(1)) or (P(2) and P(1) and G(0)) or
           |   |   | (P(2) and P(1) and P(0) and C(0));
    C(4) <= G(3) or (P(3) and G(2)) or (P(3) and P(2) and G(1)) or
           |   |   | (P(3) and P(2) and P(1) and G(0)) or
           |   |   | (P(3) and P(2) and P(1) and P(0) and C(0));

    -- Sum
    S <= P xor C(3 downto 0);

    -- Final carry-out
    C_out <= C(4);
end behavioral;
```


Subtractor

- Performed by taking the two's complement of the second number, then adding
- To compute $Y = A - B$
 - first create the two's complement of B
 - invert the bits of B to obtain \bar{B}
 - add 1 to get $-B = \bar{B} + 1$
 - add this quantity to $A -$
 - $Y = A + \bar{B} + 1 = A - B$
- We can use a single carry-lookahead adder by adding $A + B$ with $C_{in} = 1$



Comparators (1)


- Determines whether two binary numbers are equal or if one is greater or less than the other
 - receives two N-bit binary numbers A and B
 - an **equality comparator** produces a single output, indicating A is equal to B
 - a **magnitude comparator** produces one or more outputs, indicating the relative values of A and B
- The equality comparator is the simpler piece of hardware



- it checks whether the corresponding bits in each column of A and B are equal
 - using XNOR gates
- the numbers are equal if all the columns are equal

Comparators (2)



- Magnitude comparison of signed numbers is usually done by computing $A - B$ and looking at the sign (most significant bit) of the result
 - if the result is negative (i.e., the sign bit is 1), then A is less than B
 - otherwise, A is greater than or equal to B
- This comparator functions **incorrectly upon overflow**
 - overflow occurs when
 - the two inputs have different signs
 - and the sign of the subtraction result has a different sign than the A input



Arithmetic/Logical Unit (1)


- ALU **combines** a variety of **mathematical** and **logical** operations into a single unit
 - addition, subtraction, AND, and OR operations
 - **control signal** to specifies the function to perform
 - the heart of most computer systems
- The following implementation contains
 - an N-bit adder, N 2-input AND and OR gates

Control		Function
C_1	C_0	
0	0	Add
0	1	Subtract
1	0	AND
1	1	OR




Arithmetic/Logical Unit (2)

- ALU can produce **flags** to provide information about the result
- **Z (zero)**
 - all the bits of the output are 0
- **N (negative)**
 - the most significant bit of the output
- **C (carry)**
 - adder produces a carry out and the ALU is performing addition or subtraction
- **V (overflow)**
 - overflow occurs when the addition of two same signed numbers produces a result with the opposite sign
 - ALU is performing addition or subtraction ($C_1=0$)
 - A and Sum have opposite signs
 - A and B have the same sign, and the adder is performing addition ($C_0=0$)
 - or A and B have opposite sign, and the adder is performing subtraction ($C_0=1$)



Arithmetic/Logical Unit (3)



Arithmetic/Logical Unit (4)



- The ALU flags can also be used for **comparisons**:

Condition	Expression	Meaning
$A = B$	$Z = 1$	The subtraction result is zero
$A \neq B$	$Z = 0$	The subtraction result is nonzero
$A < B$	$N \text{ xor } V = 1$	Negative and overflow flags differ
$A \leq B$	$Z \text{ or } (N \text{ xor } V) = 1$	Either equal or less than
$A > B$	$Z \text{ and } (N \text{ xor } V) = 0$	Not equal and not less than
$A \geq B$	$N \text{ xor } V = 0$	Negative and overflow flags are the same


- Processors can use combinations of status bits for **conditional branching** (e.g., checking if a value is less than zero or within bounds)
- Many variations** on this basic ALU exist that support other functions

Shifters (1)

- Shifters **moves binary data**, crucial in arithmetic operations, logical operations, and data manipulation. The main types are:
 - logical shifter**: shifts all bits to the left (or right), inserting a 0 into lsb (or msb) and discarding msb (or lsb)



- arithmetic shifter**: like a logical shift with sign bit preserved if possible



- A left shift by N bits multiplies the number by 2^N
 - $000011_2 \ll 4 = 110000_2$ is equivalent to $3_{10} \times 2^4 = 48_{10}$
- A right shift by N bits multiplies the number by 2^N
 - $11100_2 \ggg 2 = 11111_2$ is equivalent to $-4_{10}/2^2 = -1_{10}$

Shifters VHDL







```
library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity Shifter is
    generic (
        N : integer := 8  -- width of the data word
    );
    port (
        A      : in  std_logic_vector(N-1 downto 0);  -- input data
        SHIFT  : in  integer range 0 to N-1;           -- shift amount
        DIR    : in  std_logic;                         -- 0 = left, 1 = right
        ARITH  : in  std_logic;                         -- 0 = logical, 1 = arithmetic
        Y      : out std_logic_vector(N-1 downto 0)     -- output data
    );
end Shifter;

architecture Behavioral of Shifter is
begin
    process (A, SHIFT, DIR, ARITH)
        variable temp : std_logic_vector(N-1 downto 0);
    begin
        if DIR = '0' then -- LEFT SHIFT
            if SHIFT < N then
                temp := std_logic_vector(shift_left(unsigned(A), SHIFT));
            else
                temp := (others => '0');
            end if;
        else -- RIGHT SHIFT
            if ARITH = '1' then -- Arithmetic right shift (preserve sign bit)
                temp := std_logic_vector(shift_right(signed(A), SHIFT));
            else -- Logical right shift (fill with zeros)
                temp := std_logic_vector(shift_right(unsigned(A), SHIFT));
            end if;
        end if;
        Y <= temp;
    end process;
end Behavioral;
```




Shifters as mux

- An N-bit shifter can be built from N N:1 multiplexers
- Depending on the value of a $\log_2(N)$ -bit shift amount control signal, the output receives the input shifted by 0 to $N-1$ bits



Rotators

- Rotator rotates (left or right) a number **in a circle** such that empty spots are filled with bits shifted off the other end




Rotators VHDL

```
entity Rotator is
    generic (
        N : integer := 8 -- width of the data word
    );
    port (
        A      : in  std_logic_vector(N-1 downto 0); -- input data
        SHIFT  : in  integer range 0 to N-1;           -- rotation amount
        DIR    : in  std_logic;                         -- 0 = left, 1 = right
        Y      : out std_logic_vector(N-1 downto 0)     -- output data
    );
end Rotator;


architecture Behavioral of Rotator is
begin
    process (A, SHIFT, DIR)
        variable temp : std_logic_vector(N-1 downto 0);
    begin
        if DIR = '0' then
            -- Rotate left
            temp := A(N-1-SHIFT downto 0) & A(N-1 downto N-SHIFT);
        else
            -- Rotate right
            temp := A(SHIFT-1 downto 0) & A(N-1 downto SHIFT);
        end if;
        Y <= temp;
    end process;
end Behavioral;
```

Multiplier (1)

- Multiplication of 1-bit numbers is equivalent to the AND operation
- A $N \times N$ multiplier produces a $2N$ -bit result
 - partial product is the AND of a single multiplier bit (B_3, B_2, \dots) with the multiplicand bits (A_3, A_2, \dots)
 - each partial products is added to the shifted next partial product
 - B_0 AND (A_3, A_2, A_1, A_0) is added to B_1 AND (A_3, A_2, A_1, A_0)




- Considering an unsigned 4×4 multiplier



Multiplier (2)

- Considering an unsigned 4×4 multiplier



Multiplier VHDL

```
entity Multiplier is
  generic (
    N : integer := 4 -- width of operands
  );
  port (
    A : in std_logic_vector(N-1 downto 0);
    B : in std_logic_vector(N-1 downto 0);
    P : out std_logic_vector(2*N-1 downto 0)
  );
end Multiplier;


architecture Behavioral of Multiplier is
begin
  process (A, B)
    variable A_u : unsigned(N-1 downto 0);
    variable B_u : unsigned(N-1 downto 0);
    variable P_u : unsigned(2*N-1 downto 0);
  begin
    -- convert to unsigned
    A_u := unsigned(A);
    B_u := unsigned(B);

    -- perform multiplication
    P_u := A_u * B_u;

    -- assign result
    P <= std_logic_vector(P_u);
  end process;
end Behavioral;
```

Division (1)

- The most complex operation to perform among all the arithmetic operations $\frac{A}{B} = Q + \frac{R}{B}$
- The basic idea is the same as long division in decimal
 - how many times the divisor can fit into sections of the dividend without exceeding it?
 - each time it fits, we add a 1 to the quotient
 - if it doesn't fit, we add a 0



- This is the **restoring division algorithm**

Division (2)

```

R' = 0
for i=N-1 to 0:
    R = R' << 1, Ai
    D = R - B
    if D<0:
        Qi = 0
        R' = R
    else:
        Qi=1
        R'=D
    R = R'

```

$A = 1101, B = 0010$

$R' = 0000$

$i = N-1 = 3$

$R = R' << 1, A_3 = 0000 << 1, 1 = 0001$
 $D = R - B = 0001 - 0010 < 0$
 $Q_3 = 0 \text{ and } R' = R = 0001$

$i = 2$


$R = R' << 1, A_2 = 0001 << 1, 1 = 0011$
 $D = R - B = 0011 - 0010 = 0001 \geq 0$
 $Q_2 = 1 \text{ and } R' = D = 0001$

$i = 1$







$R = R' << 1, A_1 = 0001 << 1, 0 = 0010$
 $D = R - B = 0010 - 0010 = 0000 \geq 0$
 $Q_1 = 1 \text{ and } R' = D = 0000$

$i = 0$

$R = R' << 1, A_0 = 0000 << 1, 1 = 0001$
 $D = R - B = 0001 - 0010 < 0$
 $Q_0 = 0 \text{ and } R' = R = 0001$



Floating-Point Addition (1)



Floating-Point Addition (2)

- Floating-point arithmetic is usually done in hardware to make it fast
 - **floating-point unit** (FPU)
 - typically, distinct from the **central processing unit** (CPU)
- The **infamous floating-point division bug** (1994)
 - the processor returns incorrect binary floating point results when dividing certain pairs of high-precision numbers...
 - missing values in a lookup table used by the FPU's division algorithm led to calculations acquiring small errors
 - these errors only occur rarely producing small deviations from the correct output
 - in certain circumstances the errors can occur frequently and lead to more significant deviations
 - **Byte magazine** estimated that **1 in 9 billion floating point divides with random parameters would produce inaccurate results**
 - **December 1994**, Intel recalled the defective processors in what was **the first full recall of a computer chip**: it cost \$475 million

